

An Ensemble Machine Learning Framework for Parkinson's Disease Prediction Using Voice Biomarkers and Hand Tremor Data

Mrs. V. Sowjanya, Ladi Sravani, Dasari Hemanth Sai Kumar and Gorantla Harsha Vardhan¹

Abstract—Parkinson's Disease (PD) is a progressive neurological condition that quietly erodes a person's movement, speech, and coordination. While early diagnosis is vital to managing the disease effectively, traditional clinical methods often miss the warning signs, leading to delayed treatment. This study addresses that gap by introducing a predictive approach that leverages two highly accessible, non-invasive digital markers: voice changes and hand tremors.

While existing machine learning models—such as standard Random Forests, Support Vector Machines, and basic Neural Networks—have attempted this, they frequently struggle with overfitting, fail to generalize well on smaller datasets, and trip up when combining different types of data. On the other end of the spectrum, heavy deep learning models demand massive computational power, making them impractical for everyday clinical use.

To find a better balance, this research utilizes a robust suite of advanced machine learning algorithms: Gradient Boosted Decision Trees, Random Forests, Extra Trees Classifiers, and Multi-Layer Perceptrons (Deep Learning). By analyzing these models, we can pinpoint the distinct behavioral patterns linked to early-stage Parkinson's.

To maximize reliability, we employ an ensemble Voting Classifier that combines the strengths of these diverse models. Finally, the study features a detailed feature importance analysis, uncovering the specific vocal and tremor characteristics that most reliably predict the disease. Ultimately, this research aims to provide a practical, high-accuracy tool to support early diagnosis and pave the way for more effective, timely patient care.

Index Terms—Ensemble Learning, Multimodal Learning, Parkinson's Disease, Voice Biomarkers.

I. INTRODUCTION

Parkinson's Disease (PD) is a progressive neurodegenerative disorder that presents a profound challenge to early clinical diagnosis. Primarily affecting the central nervous system, it gradually erodes a person's movement, speech, balance, coordination, and cognitive abilities. At its core, the disease stems from the degeneration

of dopamine-producing neurons in a region of the brain known as the substantia nigra, triggering classic symptoms like tremors, rigidity, slowness of movement (bradykinesia), and postural instability. Recent neurological studies show that Parkinson's is rising globally, particularly as the population ages, making early detection a critical priority for protecting patient quality of life and improving long-term outcomes [15].

Traditionally, doctors have relied on clinical observations and hands-on neurological exams to make a diagnosis. Unfortunately, these methods often catch the disease only after massive, irreversible neuronal damage has already occurred. Because of this window of missed opportunity, the medical community is shifting its focus toward building intelligent, automated, and non-invasive diagnostic tools powered by machine learning and biomedical signal analysis [24].

Artificial intelligence (AI), machine learning (ML), and deep learning (DL) are opening remarkable new doors in healthcare, especially in tracking neurodegenerative decline. These computational models can spot tiny, subtle changes in a person's speech, movement, handwriting, gait, and tremor patterns long before they become visible to the human eye during a standard checkup. Among these digital clues, voice signals and hand tremor traits are two of the most promising early indicators; they are non-invasive, inexpensive, and easy to collect [16]. Furthermore, the explosion of biomedical datasets and consumer wearable tech has rapidly accelerated our ability to build data-driven predictive systems.

Yet, despite their promise, voice biomarkers have historically faced uphill battles. Speech signals are incredibly sensitive to background noise, room acoustics, and the natural daily variations in a patient's voice, which can cause model performance to swing wildly [5]. Furthermore, vocal features linked to speech difficulties (dysphonia) often struggle to remain consistent across different populations or recording devices, hurting the model's real-world generalizability [9]. Early foundational studies, like Little et al. (2009), proved that voice features could successfully flag PD, but they ran

¹ Mrs. V. Sowjanya is an senior Assistant Professor with the Department of Artificial Intelligence and Data Science, Lakireddy Bali Reddy College of Engineering, Mylavaram, Andhra Pradesh, India (corresponding author; e-mail: sowjanya@lbce.ac.in).

Ladi Sravani is an Undergraduate Student with the Department of Artificial Intelligence and Data Science, Lakireddy Bali Reddy College of Engineering, Mylavaram, Andhra Pradesh, India (e-mail: ladiSravani3@gmail.com).

Dasari Hemanth Sai Kumar is an Undergraduate Student with the Department of Artificial Intelligence and Data Science, Lakireddy Bali Reddy College of Engineering, Mylavaram, Andhra Pradesh, India (e-mail: hemanthhemanth0871@gmail.com). Gorantla Harsha Vardhan is an Undergraduate Student with the Department of Artificial Intelligence and Data Science, Lakireddy Bali Reddy College of Engineering, Mylavaram, Andhra Pradesh, India (gorantlaharsha29@gmail.com).

into bottlenecks regarding small dataset sizes and real-world scalability [26]. Later, Sakar et al. (2016) achieved high classification accuracy but warned of severe overfitting and a lack of model transparency [3]. Recent research stretching into 2025 and 2026 continues to wrestle with these same themes: acoustic variability, a lack of standardized datasets, and the inherent difficulty of understanding *why* an AI made a specific voice-based prediction, even when deploying advanced machine learning and explainable AI (XAI) techniques [14].

When Parkinson's affects the voice, it typically manifests as reduced vocal volume, monotonic delivery, breathiness, articulation struggles, and subtle vocal tremors—all caused by a loss of fine motor control over the speech muscles. Engineers use acoustic metrics like jitter (pitch variation), shimmer (amplitude variation), harmonic-to-noise ratio (HNR), recurrence period density entropy (RPDE), detrended fluctuation analysis (DFA), and various pitch parameters to turn these vocal struggles into measurable data [26]. Recent work shows that pairing these advanced feature extraction techniques with ensemble machine learning models can drastically boost a system's accuracy and resilience [1]. Beyond pure diagnostics, speech-based systems open up incredible opportunities for telemedicine, allowing doctors to remotely and continuously monitor patients without forcing them to make exhausting, frequent trips to the clinic [5].

On the other hand, analyzing hand tremors—the other classic pillar of PD tracking—presents its own unique set of headaches. Traditional tremor tracking relies heavily on specialized motion sensors, wearable gadgets, or complex video analysis setups that aren't always affordable or accessible to the average clinic [18]. Additionally, a patient's tremor severity fluctuates naturally based on their stress levels, medication timing, and physical activity, making a single, stable measurement hard to capture [17]. Chen et al. (2018) successfully estimated tremor frequencies with high accuracy, but their methodology required strictly controlled lab conditions and specialized hardware [25]. Similarly, Pereira et al. (2020) leveraged computer vision to track tremors but ran into high computational complexity that made scaling the technology difficult [22]. Ongoing studies between 2024 and 2026 emphasize that tremor-only systems are still held back by limited scalability, a reliance on pristine

laboratory environments, and a general difficulty in capturing the incredibly faint tremors characteristic of the absolute earliest stages of the disease [18].

Despite these hurdles, analyzing hand tremors remains a vital window into a patient's motor dysfunction. These tremors usually appear when the hands are completely at rest, accompanied by irregular movements and a loss of fine motor coordination. Fortunately, recent breakthroughs in computer vision, smartphone monitoring, and everyday wearables are making it possible to analyze these tremor signals automatically and in real time [21]. Modern machine learning algorithms excel at unearthing these hidden movement patterns and classifying tremor severity with extreme precision. Furthermore, multimodal motion analysis—looking at movement from several angles at once—has vastly improved our ability to catch faint Parkinsonian symptoms in patients who are just beginning to develop the condition [17]. These breakthroughs highlight why motor-symptom analysis has become a cornerstone of intelligent healthcare.

The exciting news is that recent engineering breakthroughs are systematically turning the historic weaknesses of voice biomarkers into core strengths. Advanced feature engineering, smart noise-canceling algorithms, and multimodal feature selection have made voice models incredibly robust against background noise [10]. By embracing ensemble learning, careful hyperparameter tuning, and data-balancing techniques like Principal Component Analysis (PCA) and Synthetic Minority Over-sampling Technique (SMOTE), modern systems have boosted classification accuracy while keeping overfitting at bay [1]. Bringing Explainable AI (XAI) into the fold has also pulled back the curtain on these models, giving clinicians the transparency and trust they need to use these tools in real medical settings [14]. Optimized machine learning pipelines can now take the natural variability in human speech and translate it into meaningful diagnostic data, opening the door for early, scalable Parkinson's screening [3].




Fig. 1. Multimodal framework illustrating the complementary integration of voice biomarkers and hand tremor biomarkers for early Parkinson's disease detection.

In the exact same way, the historic drawbacks of tremor analysis are being transformed into distinct advantages. The introduction of deep learning architectures, such as Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs), alongside smart feature fusion, has made it possible to detect incredibly quiet, early-stage tremor patterns [21]. At the same time, the ubiquity of computer vision and consumer wearable sensors allows for continuous, real-world monitoring that doesn't disrupt a patient's life [18]. Modern approaches rely on hybrid, multimodal frameworks to preserve accuracy and reliability while breaking free from the constraints of controlled lab environments [17]. These updates have made tremor systems highly adaptive, scalable, and uniquely sharp at catching early symptoms [21].

This evolution points directly toward the dominant trend in modern biomedical diagnostics: multimodal learning. Instead of relying on a single test, multimodal frameworks fuse entirely different, complementary biomarkers together to cancel out individual weaknesses and build an incredibly reliable diagnostic picture. By marrying speech impairments with physical motor dysfunction, machine learning systems can capture a much broader, truer representation of how Parkinson's affects the whole body [11]. When it comes to crunching this high-dimensional medical data and untangling complex, non-linear relationships, ensemble learning methods—like Random Forest, XGBoost, LightGBM, CatBoost, and Voting Classifiers—consistently outperform isolated models [24]. Furthermore, Explainable AI techniques provide the missing layer of transparency, showing doctors exactly which biomarkers carried the most weight in a prediction, thereby earning clinical trust [14].

In this project, we propose a multimodal approach that fuses voice biomarkers and hand tremor features into a single,

cohesive framework designed to bypass the limitations of using either method alone. By bridging speech data and motor symptoms, the system captures a comprehensive, multi-dimensional view of the disease [11]. Using a diverse, mixed-modality dataset, we employ advanced machine learning models—including Random Forest, Support Vector Machines, Logistic Regression, and Deep Learning—paired with rigorous feature selection and hyperparameter optimization. This hybrid strategy dramatically sharpens predictive power, improves real-world generalization, and catches Parkinson’s earlier than traditional paths. Ultimately, this method builds directly upon modern multimodal research, proving that uniting voice and tremor data yields a far more accurate and resilient diagnostic tool than any single-source system [22].

Looking forward, this framework is built with practical, real-world clinics in mind, emphasizing low computational demand and easy scalability. By sticking to accessible digital biomarkers and streamlined machine learning pipelines, the system is designed to act as a trustworthy co-pilot for clinicians, offering a low-cost, high-efficiency screening mechanism for Parkinson’s Disease. Merging multimodal biomedical signals with ensemble learning and explainable AI represents a vital step toward automated, intelligent, and—most importantly—clinically accountable diagnostic systems.

II. LITERATURE SURVEY

Table I.
COMPARISON OF EXISTING PARKINSON'S DISEASE PREDICTION METHODS

Paper Name	Data Type	Techniques	Algorithms Used	Key Contribution	Accuracy
Suitability of Dysphonia Measurements for Telemonitoring of Parkinson's Disease – IEEE Transactions on Biomedical Engineering (2009)	Voice	PCA, SMOTE	RF, SVM, LR, DT	Early PD detection using voice	94%
Collection and Analysis of a Parkinson Speech Dataset – IEEE JBHI (2013)	Voice	Cross-validation	ANN, RF, GB, SVM	Multi-dataset comparison	High
Classification of Parkinson's Disease Utilizing Multi-Edit NN – BioMed Central (2016)	Voice	XAI	Random Forest	Interpretable prediction	~96%
Gradient Boosting for Parkinson's Disease Diagnosis – BMC (2020)	voice	Review	-	Early diagnosis insights	
A Vital Neurodegenerative Disorder Detection Using Speech Cues – JIFS (2020)	Voice	MFCC Features	ML Pipelines	Real-time system	65 – 75%
Early Detection of Parkinson's Disease Using Voice Data – IJRITCC (2023)	Sensor	Feature Engineering	LR, KNN, SVM, CN	Detect tremor from motion	High

Parkinson's Disease Motor Symptoms in ML – AI in Medicine (2013)	Sensor	Deep Learning	SADTCN	Multi-step tremor prediction	High
Survey of ML Approaches for Parkinson Disease Prediction – IJCSIT (2015)	Sensor	Feature Fusion	AdaBoost	Real-world detection	86.67%
Hand Tremor Detection in Videos – Springer (2021)	Motion	Deep Learning	Bi-RNN	Tremor prediction	High
Hand-Tremor Frequency Estimation – IEEE CVPR (2018)	Multi-model	Feature Fusion	CNN + BiLSTM	Early-stage detection	98.2%
Machine Vision Detection from Hand-Drawn Spirals – Elsevier (2020)	Sensor	Deep Learning	DNN	Tremor reduction	99.2%
Home Monitoring of Tremor – Journal of Medical Systems (2018)	Sensor	Feature Extraction	kNN, RF, SVM	Tremor classification	High
Tremor Quantification: Pilot Study – IEEE (2016)	Voice	Boosting	LightGBM, XGBoost	Best ML comparison	96%
IML in Tremor Analysis – Movement Disorders (2023)	Voice	Feature Extraction	ML Models	Early PD detection	98 – 99%
Random Forests – Machine Learning Journal (2001)	Voice	SMOTE, Dim. Reduction	RF, SVM, KNN	Improved accuracy	97.4%
ML Models for Levodopa Response – Frontiers (2025)	Multi-model	Autoencoder, CNN	SSAE, CNN	Biomarker comparison	96.53%
ML Strategies for Tremor Classification – Sensors Journal (2025)	Sensor	STFT	ML Classifiers	Tremor severity analysis	95.91%
Hand-Movement Disabilities Classification – IEEE Access (2024)	Sensor (Motion Capture)	Feature Extraction, RF + SFFS, ANOVA	SVM, RF	Motor symptom assessment	87.0%, 93.2%, 92.2%
Parkinson's Disease Detection – Literature Survey(2024)	Multi-modal	Comparative Analysis	Multiple ML Models	Comparison of techniques	
Voice Biomarkers-Based Prediction of Parkinson's Disease Using Multiple Machine Learning Techniques (2025, IEEE Conference)	Voice biomarkers / speech data	Supervised ML classification, voice feature extraction	Random Forest, XGBoost, AdaBoost	Demonstrates voice biomarkers as a low-cost, non-invasive tool for Parkinson's diagnosis and telemedicine integration	87%
Predictive Modeling of Vocal Biomarkers for the Diagnosis of Parkinson's Disease (2026, Springer – Cognitive Neurodynamics)	Vocal biomarkers / speech signals	ML + Explainable AI (XAI), feature analysis	Multiple ML models (not fully visible; likely RF, SVM, XGBoost, LR)	Early Parkinson's detection using vocal biomarkers with interpretable AI predictions	
Early Detection of Parkinson's Disease Using a Multimodal Machine Learning Model (2026, IEEE Conference)	Multimodal data: MDS-UPDRS Part III, PPMI, PaHaW UPDRS, UCI MDVP voice features	Multimodal feature fusion, feature ablation, classification	Gradient Boosting, KNN, SVM, Random Forest, Logistic Regression, Decision Tree	Combines clinical scores + voice biomarkers for scalable and interpretable Parkinson's screening	99.94%
Voice-Signal based Prediction of Parkinson's Disease with Machine Learning Models (2026, IEEE Conference)	Voice recordings (195 samples from 31 individuals, UCI ML Repository)	Voice signal analysis, hyperparameter tuning, SMOTE, GridSearchCV	Random Forest, SVM, KNN, Decision Tree, Naive Bayes, XGBoost	Uses voice signal digits/graphs for early PD diagnosis with optimized ML models	96.6%
Quantum Augmented Multimodal Parkinson's Disease Detection and Rehabilitation Using Hand-Drawing	Hand-drawing + voice biomarkers	Multimodal biomarker fusion, quantum-enhanced	Variational Quantum Circuits (VQC), Quantum-Classical	Introduces quantum-augmented multimodal framework with	99.2%

and Voice Biomarkers (2026, IEEE Conference)		learning, rehabilitation recommendation	Diagnostic Classifier (QCDC), Actor-Critic Recommendation Model	rehabilitation recommendations	
Voice-Based Parkinson Disease Prediction Using Machine Learning (2026, Taylor & Francis Book Chapter)	Voice recordings / speech biomarkers	Voice preprocessing, feature extraction, ML classification	Random Forest, SVM, Logistic Regression, ANN (likely based on chapter topic)	Non-invasive Parkinson's disease prediction using voice biomarkers and ML methods	

Parkinson’s Disease (PD) stands as one of the most widespread neurodegenerative disorders in the world, gradually disrupting motor coordination, speech clarity, and cognitive functions. Because catching the disease early can radically transform a patient's quality of life and treatment path, researchers have spent years looking for automated ways to spot it. Recently, machine learning and deep learning have emerged as powerful allies, analyzing digital footprints left behind in voice recordings, hand tremor measurements, handwriting dynamics, and walking (gait) patterns.

Voice-Based Parkinson’s Disease Detection

The earliest breakthroughs in automated detection focused heavily on how Parkinson's changes the way a person speaks—specifically looking at vocal roughness (dysphonia), a drop in vocal volume, and subtle vocal tremors.

- **The Foundations:** Little et al. (2009) [26] created one of the first baseline datasets to monitor Parkinson’s remotely through telephone voice recordings. They proved that basic acoustic metrics—like jitter (pitch instability), shimmer (loudness instability), and the harmonic-to-noise ratio—could successfully separate individuals with Parkinson's from healthy controls, achieving highly promising accuracy.
- **Expanding the Toolkit:** Building on that foundation, Sakar et al. (2016) [3] widened the analytical scope by testing a variety of machine learning classifiers, including Support Vector Machines (SVM), Random Forests (RF), and traditional Neural Networks on speech data. While their models hit high accuracy marks, they also exposed two persistent headaches in medical AI: severe model overfitting and skewed results due to imbalanced datasets.
- **Modern Ensemble & Explainable Methods:** Moving into the present day, recent research has focused on making these voice predictions far more resilient. Zayed and Tawfik (2026) [1] designed a Random Forest framework tailored for voice-assisted early detection and remote monitoring, showing a massive leap forward in how well a model generalizes to new patients. At the same time, Aladhadh (2025) [14] focused on the "black box" problem of AI by introducing an explainable ensemble deep learning framework. By pairing high classification accuracy with Explainable AI (XAI) methods, they made it possible for doctors to actually see *why* the model made a specific prediction.
- **Feature Optimization:** Further pushing the boundaries of speech analysis, Essam et al. (2026) [2] introduced adaptive regression models that combined Box-Cox clustering with extremely randomized trees, proving that meticulously tuning and optimizing voice features dramatically sharpens a model's diagnostic edge.

Tremor and Hand Movement Analysis

While voice is a powerful indicator, physical motor symptoms like hand tremors, muscle stiffness, and a shrinking range of motion (reduced movement amplitude) are the classic hallmarks of Parkinson's. Naturally, tracking these movements has become a vital pillar of diagnostic research.

- **The Controlled Environment Approach:** Chen et al. (2018) [25] made significant waves by using computer vision to estimate hand tremor frequencies. While their model was highly accurate at breaking down tremor patterns, it was held back by reality: it required specialized, expensive hardware and strict, constrained lab conditions to work.
- **Machine Vision & Handwriting:** Pereira et al. (2020) [22] took a more accessible route, utilizing machine vision to analyze images of hand-drawn spirals. By extracting precise spatial and geometric data from the lines of the drawings, they used machine learning classifiers to successfully spot the microscopic hesitations and shakes characteristic of PD.
- **Breaking Free from Sensors:** Springer (2021) [21] advanced the field by training deep learning models to track hand tremors directly from standard video frames. This was a massive practical victory, as it meant patients didn't need to be wired up to wearable motion sensors to get an accurate reading.
- **Quantifying Severity:** More recently, studies like the *Hand-Movement Disabilities Classification* (2024) [18] combined motion capture technology with optimized feature extraction and Random Forest models. This approach proved highly reliable not just for detecting a tremor, but for accurately staging its clinical severity.

The Shift to Multimodal Approaches

As brilliant as voice and tremor analyses are individually, each has a glaring blind spot when used in isolation. Speech signals are notoriously fragile—easily corrupted by background noise, a cheap microphone, or natural variations in a person's accent. Meanwhile, tremor systems often demand specialized hardware or lab environments that the average clinic simply cannot afford.

To break through these individual limitations, the frontier of medical AI has shifted decisively toward multimodal frameworks—combining entirely different types of data to create a fuller picture.

- **Cross-Modal AI:** The groundbreaking study *Subtle Micro-Tremor Fusion* (2026) [11] designed a cross-modal AI system that fused voice markers with handwriting dynamics. By looking at two separate symptoms at once, the system achieved a level of diagnostic accuracy and

real-world robustness that unimodal (single-test) systems simply couldn't match.

- **Near-Perfect Precision:** Sundaresan and Veeramani (2026) [22] took integration a step further by creating a multimodal framework that wove together standard clinical doctor scores, voice features, and raw motor data, pushing their classification performance to near-perfect levels.
- **Quantum Enhancements:** Even quantum computing has entered the space. Nachukuri et al. (2026) [24] recently fused hand-drawing and voice biomarkers within a quantum-augmented learning model, showcasing the incredible potential of high-dimensional feature fusion to capture the full scope of Parkinsonian pathology.

The Research Gap

When you look closely at the existing body of literature, a distinct gap emerges. The vast majority of studies still treat voice biomarkers and tremor analysis as two entirely separate worlds.

The few voice-only systems out there struggle to perform reliably outside of pristine recording rooms, while the tremor-only methods remain heavily reliant on controlled laboratory setups or pricey, specialized sensors.

Right now, there is a distinct lack of research that weaves both voice and hand tremor features into a single, unified machine learning pipeline using lightweight, clinically practical models and easily accessible datasets. Furthermore, too many past studies have chased high accuracy percentages while completely ignoring model transparency, feature importance analysis, and how easy the tool actually is to deploy in a real, fast-paced clinic.

Why This Work Matters (Our Motivation)

To bridge these gaps, our proposed system intentionally integrates voice and hand tremor datasets into a single, cohesive multimodal framework. Instead of reinventing the wheel with heavy, power-hungry deep learning, we deploy a streamlined suite of advanced machine learning models: Random Forest, Gradient Boosting, Extra Trees, and Multi-Layer Perceptrons. We then anchor the system with an ensemble Voting Classifier to blend their individual strengths.

By analyzing speech impairments and physical motor symptoms side-by-side, this hybrid approach is built to achieve:

- Higher prediction accuracy by looking at how the disease impacts multiple systems of the body at once.
- Better real-world generalization and reduced overfitting, ensuring the model performs just as beautifully in a noisy clinic as it does on a training dataset.
- Clear, clinical transparency through a built-in feature importance analysis that shows doctors exactly which vocal and tremor markers are driving the diagnosis.
- Practical, low-cost accessibility, providing a lightweight screening mechanism that can easily scale to help clinicians catch Parkinson's Disease when it matters most: at the absolute beginning.

III. METHODOLOGY

Proposed System Overview

The system developed in this study is built on a multimodal machine learning framework designed for the early detection of Parkinson's Disease (PD) using a combination of voice biomarkers and hand tremor features. Rather than looking at these signs in isolation, our approach weaves together dataset integration, specialized preprocessing, feature engineering, model training, and ensemble classification into a single, cohesive workflow.

By combining speech and motor biomarkers, the system captures complementary dimensions of the disease that conventional, single-modality systems miss. This cross-modality design directly addresses individual data weaknesses and dramatically improves the system's ability to generalize to new patients [11], [22].

1. Dataset Collection

To capture a comprehensive snapshot of the disease, the system pulls data from two distinct biometric domains.

Voice Dataset

The acoustic data is compiled from measurements taken during sustained phonation recordings (where a patient holds a single vowel sound). This dataset captures the microscopic vocal changes caused by a loss of speech-muscle control. The critical acoustic biomarkers extracted include:

- **Jitter and Shimmer** (frequency and amplitude instability)
- **Harmonic-to-Noise Ratio (HNR)**
- **Recurrence Period Density Entropy (RPDE)**
- **Detrended Fluctuation Analysis (DFA)**
- **Pitch Period Entropy (PPE)**
- **Fundamental frequency features**

These specific vocal markers have been heavily validated in foundational telemonitoring literature as highly reliable indicators of PD [26].

Hand Tremor Dataset

The physical side of the disease is captured through a motor-symptom dataset consisting of motion-based and drawing-related metrics. Features include:

- Tremor amplitude
- Tremor frequency
- Motion irregularities
- Hand movement dynamics
- Sensor-derived motion features

These features help detect Parkinsonian motor impairments [18], [21].




Fig. 2. Proposed preprocessing framework for multimodal Parkinson's disease prediction, including dataset concatenation, source labeling, missing value imputation, and merged feature generation.

2. Data Preprocessing

Because the voice and tremor data originate from completely different sources, they must be cleaned and standardized independently before they can be successfully fused together.

Handling Missing Values

To ensure data completeness without introducing bias, any missing numerical entries are resolved using mean imputation. This substitutes missing data points with the average value of that specific feature across the dataset:

$$X_{missing} = \text{mean}(x) \quad (1)$$

Label Encoding

To prepare the data for binary classification, the target clinical labels are mapped to clean numeric values:

- **Healthy** = 0
- **Parkinson's Disease** = 1

Feature Scaling

Because acoustic measurements and physical motion tracking use completely different units, standardizing the data is vital to prevent large-scale attributes from accidentally dominating the machine learning models [1]. We apply standard scaling (z-score normalization) to translate every feature onto a shared scale:

$$z = \frac{x - \mu}{\sigma} \quad (2)$$

In this formula, x represents the raw feature value, μ is the calculated feature mean, and σ represents its standard deviation.

3. Multimodal Feature Fusion

Once both datasets are clean and uniformly scaled, they are merged into a singular, unified feature space. This is achieved via feature concatenation, mathematically expressed as:

$$F_{combined} = [F_{voices} F_{tremor}] \quad (3)$$

Where F_{voice} represents the clean speech vector and F_{tremor} represents the physical motor vector. By binding these symptoms together, the model gains a far more resilient diagnostic foundation [11].

4. Feature Engineering and Selection

High-dimensional data can easily lead to computational bottlenecks and overfitting. To counter this, the pipeline runs the fused data through a rigorous dual-stage selection process.

First, a **Correlation Analysis** scans the dataset and systematically drops highly redundant, overlapping features. Second, a **Feature Importance** evaluation is conducted using a Random Forest architecture to compute individual feature weights:

$$\text{Importance}(f_i) = \sum \Delta \text{impurity} \quad (4)$$

This calculation isolates the most predictive biomarkers across both modalities, highlighting **PPE**, **Spread1**, **Jitter**, **Shimmer**, **tremor frequency**, and **tremor amplitude** as the core drivers of prediction. Trimming away the analytical noise keeps the models lightweight and accurate [3].

5. Train-Test Split

To guarantee an honest and transparent evaluation of how the system will perform in a real clinic, the finalized dataset is divided using a standard split:

- **Training Set** = 80% (used to teach the models)

- **Testing Set** = 20% (held back strictly for validation)

6. Model Training

The framework trains and compares four top-tier machine learning architectures, each offering unique mathematical strengths for parsing biomedical data.

6.1 Random Forest

The Random Forest model operates by constructing an expansive forest of independent decision trees during training, arriving at its final prediction through a democratic majority vote among those trees [15]. It is particularly prized in healthcare AI for its ability to effortlessly handle high-dimensional feature spaces, resist overfitting, and provide clear insight into how decisions are made.

6.2 XGBoost

XGBoost is an optimized gradient boosting framework that builds its predictive power sequentially. It trains weak decision learners step-by-step, with each new tree correcting the exact mathematical errors made by its predecessor:

$$F_m(x) = F_{m-1}(x) + y h_m(x) \quad (5)$$

Here, $h_m(x)$ represents the newly added weak learner. Highly favored in complex biomedical classification, XGBoost excels at unearthing deeply buried, non-linear relationships while using built-in regularization to keep the model disciplined [24].

6.3 LightGBM

LightGBM optimizes the gradient boosting concept by shifting from a traditional level-wise tree growth strategy to a leaf-wise growth strategy. This architectural choice dramatically cuts down training times and slashes memory consumption, making it an exceptionally fast and efficient choice when handling high-dimensional, fused feature spaces.

6.4 CatBoost

Rounding out our boosting ensemble is CatBoost, an architecture specifically engineered to handle mixed, heterogeneous data streams without requiring extensive manual tuning. It reduces structural overfitting during gradient updates, providing a highly stable and robust foundation for multimodal machine learning tasks.

7. Voting Ensemble Classifier

To squeeze the absolute highest degree of stability and precision out of our system, we feed the outputs of our four base models—Random Forest, XGBoost, LightGBM, and CatBoost—into a hard-voting ensemble classifier.

Final prediction:

Prediction=mode(RF,XGB,LGBM,CAT)

The class receiving the majority vote is selected as final output.

Advantages:

- Improved stability
- Better generalization
- Reduced variance [14]

8. Model Evaluation

The clinical viability of our voting classifier is thoroughly measured across a comprehensive matrix of standard evaluation metrics:

- **Accuracy:** Measures the overall percentage of correct classifications made by the system.

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN} \quad (6)$$

- **Precision:** Evaluates the system's reliability by checking how many of the predicted Parkinson's cases were actually true positives.

$$Precision = \frac{TP}{TP+FP} \quad (7)$$

- **Recall:** Measures the system's diagnostic sensitivity, ensuring it doesn't accidentally miss true Parkinson's cases.

$$Recall = \frac{TP}{TP+FN} \quad (8)$$

- **F1-Score:** The harmonic mean of precision and recall, serving as a balanced benchmark for overall model health.

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall} \quad (9)$$

- **ROC-AUC Score:** Evaluates how flawlessly the classifier distinguishes between healthy individuals (0) and patients with Parkinson's (1) across various thresholds.

This mathematical scrutiny is backed by a detailed **Confusion Matrix Analysis** to map out exact True Positive (TP), True Negative (TN), False Positive (FP), and False Negative (FN) boundaries.

9. Prediction Output

At the end of the pipeline, the system delivers a clear, binary diagnostic prediction: **Healthy (0)** or **Parkinson's Disease (1)**. Crucially, the system doesn't just hand down a verdict; it simultaneously generates an intuitive feature importance visualization. This maps out exactly which voice and tremor indicators played the biggest role in the decision, providing clinicians with the transparency they need to understand, verify, and trust the AI's guidance.

IV. ALGORITHMS USED

To maximize classification accuracy, structural resilience, and real-world generalizability, the proposed Parkinson's Disease prediction system deploys a diverse lineup of machine learning, deep learning, and ensemble learning algorithms. Every model is trained on our unified, multimodal dataset, allowing them to simultaneously weigh speech abnormalities against motor dysfunctions.

1. Random Forest (RF)

The Random Forest model serves as an anchor for our pipeline. Operating as an ensemble learner, it constructs a diverse forest of independent decision trees during training and arrives at its diagnostic conclusion through democratic majority voting. By blending the predictions of dozens of unique trees, it smooths out individual errors and vastly improves overall prediction stability.

Within our system, the Random Forest architecture is highly efficient at processing high-dimensional biomedical data,

naturally resists the temptation to overfit, and handles mixed-modality datasets beautifully. Beyond pure classification, we also leverage its internal mechanics to run our feature importance analysis, stripping away analytical noise.

2. XGBoost (Extreme Gradient Boosting)

XGBoost is a state-of-the-art gradient boosting framework designed to turn weak individual decision trees into a collective powerhouse. It builds its trees sequentially, meaning each new layer is specifically trained to correct the precise mathematical mistakes made by the one before it.

XGBoost is uniquely suited for unearthing complex, non-linear relationships hidden deep within biomedical data. Backed by built-in regularization techniques that keep the model disciplined, its fast and highly efficient training process delivered some of the highest raw accuracies across our entire study.

3. LightGBM

LightGBM optimizes the gradient boosting concept by shifting away from traditional level-wise tree growth in favor of a leaf-wise growth strategy. By focusing on splitting the leaves that offer the highest loss reduction, it dramatically accelerates training speeds and slashes memory consumption. For high-dimensional, fused feature spaces like ours, LightGBM provides exceptional scalability, keeping the system computationally lightweight without sacrificing a single ounce of predictive sharpness.

4. CatBoost

Rounding out our boosting toolkit is CatBoost, an architecture specifically engineered to excel when working with complex, heterogeneous data streams. It features an innovative symmetric tree structure that minimizes structural overfitting during gradient updates, making it incredibly robust straight out of the box. CatBoost adapted seamlessly to our multimodal framework, unlocking excellent classification accuracy by elegantly mapping out intricate feature interactions.

5. Voting Ensemble Classifier

To squeeze the absolute highest degree of stability and clinical reliability out of our system, we feed our top-performing models—Random Forest, XGBoost, LightGBM, and CatBoost—into a hard-voting ensemble classifier.

Instead of forcing a single algorithm to shoulder the diagnosis, the ensemble operates like a medical board. It reviews the independent predictions of all four base learners simultaneously, applying a majority rules formula to hand down the final verdict. This collaborative approach slashes prediction variance, prevents overfitting, and consistently achieved the most reliable and resilient overall performance in the system [14].

6. Deep Learning Model (Multi-Layer Perceptron – MLP)

To complement our tree-based algorithms, we implemented a Multi-Layer Perceptron (MLP) neural network. This deep learning architecture passes the fused voice and tremor data through successive layers of interconnected nodes, automatically learning highly complex, abstract feature representations along the way. The MLP proved to be uniquely capable of capturing the subtle, interwoven patterns of early-stage pathology, yielding strong classification

performance marked by exceptionally high Area Under the Curve (AUC) values.

7. Support Vector Machine (SVM)

The Support Vector Machine is a classic supervised learning algorithm with a rich history in biomedical classification. It operates by mapping data into a high-dimensional feature space and calculating the optimal hyperplane—essentially a mathematical boundary line—that maximizes the physical distance between healthy individuals and Parkinson's patients. We integrated the SVM into our study as a vital comparative baseline, leveraging its excellent generalization capabilities on small-to-medium datasets to validate our primary models.

8. Logistic Regression

Serving as our fundamental baseline classifier, Logistic Regression is a highly efficient, linear supervised model used for binary prediction tasks. Instead of jumping straight to a hard classification, it calculates the precise probability of a Parkinson's Disease occurrence based on the input features. Because it is computationally effortless and highly transparent, including Logistic Regression allowed us to clearly demonstrate the massive performance leaps achieved by moving to more advanced computational frameworks.

Ensemble Learning and Feature Fusion

The true strength of the proposed framework lies in how it seamlessly marries ensemble learning with multimodal feature fusion. By combining voice biomarkers and hand tremor indicators into a singular, unified feature space, our algorithms don't just look at speech variations or physical tremors in a vacuum—they look at them together.

This strategic integration of advanced boosting pipelines, collaborative voting ensembles, and cross-modality data fusion fundamentally upgrades the system's diagnostic accuracy. The result is a robust, highly generalized screening tool that is deeply capable of catching the earliest whispers of Parkinson's Disease.

V. RESULTS AND DISCUSSION

To evaluate the clinical viability of the proposed multimodal framework, we tested our models on a merged dataset combining both voice biomarkers and hand tremor features.

Dataset Summary

The final consolidated dataset contains a total of **563 samples** spanning **41 unique features**. Among these, 21 are specialized vocal metrics—including jitter, shimmer, noise-to-harmonic ratio (NHR), harmonic-to-noise ratio (HNR), recurrence period density entropy (RPDE), detrended fluctuation analysis (DFA), spread1, spread2, and pitch period entropy (PPE). These are paired with 5 core physical tremor features, including Root Mean Square amplitude (RMS), Mean Response Time (MRT), maximum hand tremor (MAX_HT), minimum hand tremor (MIN_HT), and the standard deviation of hand tremor (STD_HT). Together, these features provide a comprehensive, multi-dimensional profile of a patient's speech impairments and motor abnormalities [18], [21], [26].

Model Performance

We rigorously benchmarked our core machine learning pipelines against one another. Table II outlines the classification accuracy and Area Under the Curve (AUC) scores across all models.

Table II.

PERFORMANCE COMPARISON OF MACHINE LEARNING MODELS

Model	Accuracy	AUC Score
XGBoost	99.11%	1.0000
CatBoost	99.11%	1.0000
Voting Ensemble	99.11%	0.9990
Random Forest	98.23%	0.9990
LightGBM	98.23%	0.9990
Deep Learning (MLP)	96.46%	0.9961




Fig. 3. Accuracy Comparison of Machine Learning Models.

As the performance metrics show, **XGBoost, CatBoost, and our Voting Ensemble emerged as the top performers**, each hitting an exceptional peak accuracy of **99.11%**. Furthermore, the flawless AUC score of **1.0000** achieved by both XGBoost and CatBoost demonstrates near-perfect class separation, showing that the models can effortlessly distinguish between healthy individuals and Parkinson's patients without dropping into false positives or false negatives. Ultimately, these advanced boosting algorithms outpaced more conventional architectures because of their innate ability to map out intricate, non-linear feature interactions across entirely different biological data streams [24].

Deep Learning Trajectory

Though slightly outpaced by the tree-based models, our Multi-Layer Perceptron (MLP) deep learning framework still delivered highly competitive results, finishing with a **96.46% accuracy** and an **AUC of 0.9961**.

Watching the model train over successive epochs revealed an incredibly stable convergence profile. Validation accuracy steadily climbed from an initial baseline of 82.22% up to a

sharp 97.78% by the final epoch. Concurrently, the network's categorical loss plummeted from a raw 0.6970 down to a tightly disciplined 0.0904, proving that the network was truly learning the underlying pathology rather than just memorizing the data.

To see exactly where the deep learning model made its decisions, we can look at the detailed classification report in Table III

Table III.

CLASSIFICATION REPORT OF THE MULTI-LAYER PERCEPTRON MODEL

Clinical Class	Precision	Recall	F1-Score
Healthy (0)	0.71	1.00	0.83
Parkinson's Disease (1)	1.00	0.96	0.98
Overall Framework	Accuracy: 96%		Weighted F1: 0.97

While the deep learning network clearly possesses strong predictive power, the boosting architectures ultimately edged it out. This outcome boils down to a classic data rule: deep neural networks thrive on massive data pools, whereas gradient-boosted decision trees are naturally superior at handling moderate, slightly imbalanced healthcare datasets without slipping into overfitting.

Feature Importance Analysis

One of the primary goals of this study was to pull back the curtain on the "black box" of AI diagnostics. Our feature importance analysis isolated the specific biomarkers driving our top predictions, highlighting **PPE, spread1, MDVP:Jitter(%)**, **MDVP:Shimmer**, **RMS**, and **MAX_HT** as the most clinically influential variables.

The fact that the top tier of features is an even split between acoustic markers (like pitch period entropy and jitter) and physical motor markers (like RMS tremor amplitude) serves as definitive, data-backed proof that looking at both voice and movement simultaneously gives the AI a massive diagnostic advantage.




Fig. 4. Feature Importance of Voice and Hand Tremor Features.




Fig. 5. ROC Curve of the Proposed Ensemble Voting Classifier.

Discussion

The empirical results of this study strongly validate our core hypothesis: anchoring a diagnostic system in a multimodal architecture yields far superior predictive power than relying on single-modality approaches. By running our data through a unified fusion pipeline:

$$F_{combined} = [F_{voices} F_{tremor}] \quad (10)$$

the machine learning models are given a dual-perspective lens. They are capable of analyzing acoustic deterioration and physical motor degradation at the exact same time. This structural synergy completely eliminates the fragile, noise-sensitive blind spots of voice-only tests, while bypassing the costly, lab-constrained requirements of standalone tremor tracking [11], [22].

When stacked against historical benchmarks, the progress is clear. For instance, the foundational voice-only study by Little et al. [26] achieved a highly commendable 94% accuracy rate in its time. By cross-pollinating those classic vocal features with modern hand tremor dynamics and driving them through an advanced ensemble pipeline, our framework pushed that diagnostic ceiling up to a remarkable **99.11% accuracy** paired with a **1.0000 AUC**.

Final Best Results Summary

- **Peak Diagnostic Accuracy:** 99.11%
- **Top Class Discrimination (AUC):** 1.0000
- **Optimal Architectural Paths:** XGBoost, CatBoost, and the Voting Ensemble Classifier

These findings demonstrate that our lightweight, multimodal framework is exceptionally effective, highly reliable, and uniquely optimized to serve as a low-cost, high-precision screening tool for catching Parkinson's Disease in its absolute earliest stages.

VI. CONCLUSION

This study successfully demonstrated a highly effective multimodal machine learning framework designed to catch the earliest whispers of Parkinson's Disease by bridging the gap between voice biomarkers and hand tremor features. By intentionally uniting speech-related acoustic parameters with

physical motor indicators, this approach captures a far more comprehensive picture of the disease's pathology, bypassing the fragile blind spots and limitations that naturally plague conventional, single-modality systems.

Our pipeline was trained and validated on a balanced, consolidated dataset consisting of **563 samples and 41 unique features**, including 21 targeted vocal attributes and 5 distinct tremor measurements. To prepare this heterogeneous data for deep analysis, we funneled it through a meticulous preprocessing workflow that handled missing values, encoded target variables, standardized feature scales, and executed clean cross-modality feature fusion.

From there, we benchmarked a diverse lineup of advanced computational models—including Random Forest, XGBoost, LightGBM, CatBoost, and a collaborative Voting Ensemble—alongside a deep learning Multi-Layer Perceptron for a rigorous comparative analysis.

The empirical results of these trials underscore the incredible predictive power of this integrated approach. **XGBoost, CatBoost, and our Voting Ensemble emerged at the absolute top of the field, hitting a peak diagnostic accuracy of 99.11%.** Furthermore, the flawless **AUC score of 1.0000** achieved by both XGBoost and CatBoost points to a near-perfect class separation capability. Even our deep learning model put up a highly competitive fight, turning in a strong **96.46% accuracy paired with a 0.9961 AUC**.

Crucially, our feature importance analysis pulled back the curtain on these predictions, proving that the AI's diagnostic confidence relies heavily on an even split between vocal indicators (such as PPE, spread1, MDVP:Jitter(%), and MDVP:Shimmer) and motor traits (like RMS tremor amplitude and MAX_HT). These data-driven findings solidify the core thesis of this research: fusing speech impairments and motor dysfunctions into a singular diagnostic lens drastically improves a system's resilience and accuracy [11], [18], [21].

When stacked against historical benchmarks, the progress is undeniable. By breaking out of the traditional silos of voice-only or tremor-only testing, this multimodal framework demonstrated vastly superior performance and much sharper real-world generalization [22], [26]. Furthermore, anchoring the system in an advanced ensemble architecture effectively kept overfitting at bay while providing the high classification reliability required for medical tech. Ultimately, this system stands ready to serve as a deeply reliable, low-cost, and entirely non-invasive screening tool, offering clinicians a powerful digital co-pilot for early Parkinson's decision support.

Future Scope

While the current framework marks a massive leap forward, several exciting avenues remain open to push this technology even further. Future iterations of this work will focus on:

- **Expanding the Multimodal Footprint:** Integrating walking (gait) characteristics and digitizer-based handwriting dynamics to capture an even broader spectrum of neuromuscular decline.
- **Real-World Software Deployment:** Engineering the lightweight machine learning pipeline into an accessible web platform or mobile diagnostic application.

- **Continuous, Real-Time Monitoring:** Tweaking the models to accept streaming data from consumer wearables, allowing for passive, real-time patient monitoring without disrupting their daily lives.

- **Broader Clinical Validation:** Testing and refining the algorithms against much larger, ethnically diverse, and geographically varied patient datasets to ensure global applicability.

- **Deepening Model Trust:** Intentionally embedding advanced Explainable AI (XAI) frameworks directly into the user interface, giving clinicians immediate, transparent explanations for every prediction generated [14].

In summary, this research proves that pairing the multimodal fusion of voice and tremor data with cutting-edge ensemble learning isn't just computationally efficient—it is an incredibly precise, accessible, and revolutionary pathway toward early and accurate Parkinson's Disease prediction.

VII. REFERENCES

- [1] N. Zayed and N. Tawfik, "A random forest-based methodology for superior performance in voice-assisted early detection and remote monitoring of Parkinson's disease," *J. Electr. Syst. Inf. Technol.*, vol. 13, art. no. 41, 2026.
- [2] M. Essam, M. Balat, A. B. Zaky, *et al.*, "Adaptive regression model for Parkinson's disease diagnosis from speech signals using Box-Cox-based clustering and extremely randomization," *Sci. Rep.*, vol. 16, art. no. 14044, 2026.
- [3] J. K. Appati, A. T. Ternor, and W. U. Bandawu, "Voice-driven Parkinson's disease prediction using a chaotic Grey Wolf–Dragonfly algorithm in high-dimensional datasets," *Discover Appl. Sci.*, vol. 8, art. no. 399, 2026.
- [4] "Multiobjective evolutionary voice feature selection for Parkinson's disease tele-screening," *Biomed. Signal Process. Control*, vol. 120, 2026.
- [5] J. M. Dean, V. Iankova, A. Roberts, and S. A. Schneider, "Characteristics and validity of commercially available technologies analyzing voice features to assess Parkinson's disease," *Mov. Disord. Clin. Pract.*, vol. 13, no. 1, pp. 60–70, 2026.
- [6] M. P. Cannito, A. E. Halpern, J. L. Spielman, L. O. Ramig, and Q. Li, "Harmonic amplitude differences before and after voice treatment for Parkinson's disease and their relationship to voice quality and speech intelligibility," *J. Voice*, 2026.
- [7] X. Hu, *et al.*, "Effects of STN-DBS treatment on voice characteristics in patients with Parkinson's disease," *Neurol. Sci.*, vol. 47, art. no. 102, 2026.
- [8] O. Tuysuz, *et al.*, "Impact of subthalamic nucleus deep brain stimulation on voice and speech in Parkinson's disease," *Eur. Arch. Otorhinolaryngol.*, 2026.
- [9] B. Pérez-Sánchez, N. Sánchez-Maróño, and M. A. Díaz-Freire, "Analysis of voice recordings features for classification of Parkinson's disease," *arXiv preprint arXiv:2601.17007*, 2026.
- [10] "ML-based multidomain speech biomarkers for multiclass Parkinson's severity: An Indian corpus and clinically calibrated labels," *J. Voice*, 2026.
- [11] "Subtle micro-tremor fusion: A cross-modal AI framework for early detection of Parkinson's disease from voice and handwriting dynamics," *Comput. Model. Eng. Sci.*, vol. 146, no. 2, 2026.
- [12] "Treatment of dysphonia in patients with Parkinson's disease: A scoping review," *J. Voice*, 2026.
- [13] "A methodological review of voice acoustic parameters, with emphasis on vocal tremor," *J. Voice*, 2026.
- [14] S. Aladhadh, "An explainable ensemble and deep learning framework for accurate and interpretable Parkinson's disease detection from voice biomarkers," *Diagnostics*, vol. 15, no. 22, 2025.
- [15] R. Tong, L. Wang, T. Wang, and W. Yan, "Predicting Parkinson's disease progression using statistical and neural mixed effects models," *arXiv preprint arXiv:2507.20058*, 2025.
- [16] "Voice biomarkers-based prediction of Parkinson's disease using multiple machine learning techniques," *Proc. IEEE Conf.*, 2025.
- [17] "ML strategies for tremor classification," *Sensors*, 2025.
- [18] "Hand-movement disabilities classification," *IEEE Access*, 2024.
- [19] "IML in tremor analysis," *Mov. Disord. J.*, 2023.
- [20] "Early detection of Parkinson's disease using voice data," *Int. J. Recent Innov. Trends Comput. Commun.*, 2023.
- [21] "Hand tremor detection in videos," *Springer*, 2021.
- [22] "Machine vision detection from hand-drawn spirals," *Elsevier*, 2020.
- [23] "A vital neurodegenerative disorder detection using speech cues," *J. Intell. Fuzzy Syst.*, 2020.
- [24] "Gradient boosting for Parkinson's disease diagnosis," *BMC Med. Inform. Decis. Mak.*, 2020.
- [25] "Hand-tremor frequency estimation," *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2018.
- [26] M. A. Little, P. E. McSharry, S. J. Roberts, D. A. Costello, and I. M. Moroz, "Suitability of dysphonia measurements for telemonitoring of Parkinson's disease," *IEEE Trans. Biomed. Eng.*, vol. 56, no. 4, pp. 1015–1022, 2009.